

Information Security Lab

Autumn Semester 2022

Module 1, Week 2 – Cryptanalysis of ECDSA

Kenny Paterson (@kennyog)

Applied Cryptography Group

<https://appliedcrypto.ethz.ch/>

Overview of today's lectures

- ECDSA recap from last week
- Breaking ECDSA with a single known nonce and with repeated nonces
- Breaking ECDSA with partially known nonces
- Introducing lattices and lattice reduction
- From breaking ECDSA to CVP
- From CVP to SVP via Kannan embedding
- Putting it all together

ECDSA Recap

ECDSA Recap

Parameters: (E, p, n, q, h, P, H) defining a curve E over field F_p with $n = q \cdot h$ points, subgroup of prime order q and generator P of order q ; H is a hash function, e.g. SHA-256 (here we assume output of H is at least bit-size of q).

KeyGen:

Set $Q = [x]P$ where x is uniformly random from $\{1, \dots, q-1\}$.

Output verification key: Q ; signing key: x .

Sign: Inputs (d, m) // d is private key; m is the message to be signed

$h = \text{bits2int}(H(m)) \bmod q$. // take $\text{len}(q)$ MSBs of $H(m)$, cast to `BigInt`, reduce mod q .

Do:

1. Select k uniformly at random from $\{1, \dots, q-1\}$. // k is called the *nonce*
2. Compute $r = \text{x-coord}([k]P) \bmod q$. // $[k]P$ is a point on E ; its x-coord is in F_p ; we consider that as an integer and reduce mod q .
3. Compute $s = k^{-1}(h + xr) \bmod q$.

Until $r \neq 0$ and $s \neq 0$. // works first try w.h.p.

Output (r, s) .

ECDSA Recap

Verify: Inputs $(Q, m, (r, s))$ // Q is verification key; m is message; (r, s) is claimed signature.

1. check that $1 \leq r \leq q-1$ and $1 \leq s \leq q-1$.
2. compute $w = s^{-1} \bmod q$.
3. compute $h = \text{bits2int}(H(m)) \bmod q$.
4. compute $u_1 = w \cdot h \bmod q$ and $u_2 = w \cdot r \bmod q$.
5. compute $Z = [u_1]P + [u_2]Q$.
6. If $(\text{x-coord}(Z) \bmod q == r)$ then output 1 else output 0.

Correctness:

Suppose (r, s) is a signature for message m under key Q . Then:

$$Z = [u_1]P + [u_2]Q = [s^{-1}h]P + [s^{-1}r]Q = [s^{-1}(h + xr)]P = [k]P.$$

Here we used $s = k^{-1}(h + xr) \bmod q$ from the signing algorithm to obtain $s^{-1}(h + xr) = k \bmod q$.

Recalling that $r = \text{x-coord}([k]P) \bmod q$ completes the argument.

Simple Attacks on ECDSA

Breaking ECDSA with A Single Known Nonce

- Suppose we have an ECDSA signature (r,s) for message m in which the nonce k is known in its entirety.

- Recall the signing equation:

$$s = k^{-1}(h + xr) \bmod q.$$

- Rearranging gives:

$$x = r^{-1}(ks - h) \bmod q.$$

- Here, $h = \text{bits2int}(H(m)) \bmod q$, so depends only on m .
- Hence recovering x , the signing key, is trivial.
- You will implement this attack as a warm-up in the lab.

Breaking ECDSA with Repeated Nonces

- Suppose we have two ECDSA signatures (r_1, s_1) for message m_1 and (r_2, s_2) for message m_2 .
- Suppose we also know that the same value k was used in both cases.
- For example, maybe the developer did not realise that k needs to be a fresh, random value for each use, or used a broken RNG.
- Sony Playstation 3 fail: <https://www.bbc.com/news/technology-12116051>
- Rearranging the signing equations gives the pair of linear equations:
$$xr_1 - ks_1 = h_1 \bmod q \quad \text{and} \quad xr_2 - ks_2 = h_2 \bmod q$$

in unknowns x and k .
- Solve using linear algebra mod q to recover both unknowns.
- Explicitly: $x = (h_1s_2 - h_2s_1) \cdot (r_2s_1 - r_1s_2)^{-1} \bmod q$.
- You will implement this attack as a second warm-up in the lab.

Breaking ECDSA with Partially Known Nonces

Breaking ECDSA with Partially Known Nonces

- The preceding examples show that the security of ECDSA is very sensitive to how nonces k are chosen.
- How sensitive exactly?
- What if the attacker could learn just a few bits of k ?
- Such information might be available via a side-channel attack.
- Examples:
 - Brumley and Taveri, "Remote Timing Attacks are Still Practical", ESORICS 2011 and <https://eprint.iacr.org/2011/232.pdf>
 - Moghimi et al., "TPM-FAIL: TPM Meets Lattice and Timing Attacks", USENIX 2020 <https://www.usenix.org/conference/usenixsecurity20/presentation/moghimi-tpm>
 - Both papers observe leakage of cases when MSBs of k are zero due to faster execution of $[k]P$ during the signing algorithm: a timing side-channel observable by a "remote" attacker.
 - Other, recent work observes partial leakage of k via cache-based side-channel attacks (a local attacker model).

Breaking ECDSA with Partially Known Nonces

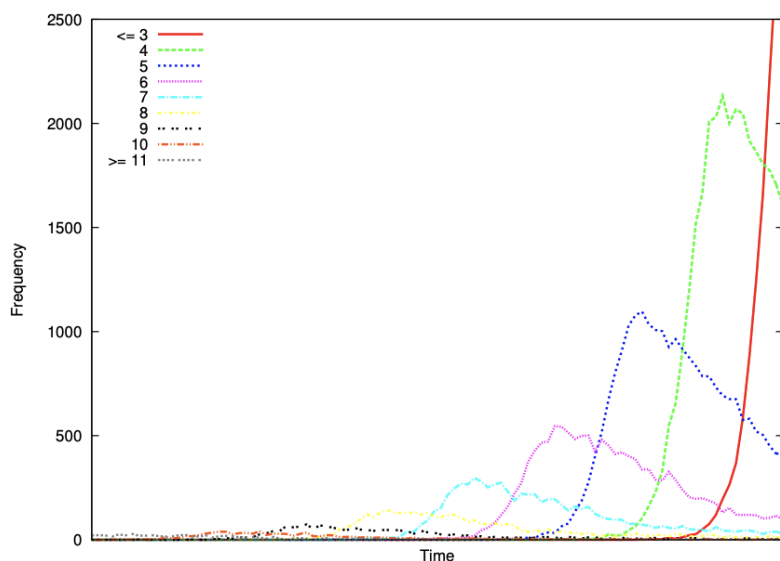


Fig. 4. Dependency between number of leading zero bits and wall clock execution time of the signature operation.

Brumley-Tuveri, 2011

OpenSSL ECDSA signature generation




Figure 2: Histogram of ECDSA (NIST-256p) signature generation timings on the STMicronics TPM as measured on a Core i7-8650U machine for 40,000 observations.

Moghimi et al., 2020

STMicronics TPM ECDSA
signature generation

Breaking ECDSA with Partially Known Nonces

Table 2: Discovered vulnerabilities in OpenSSL, LibreSSL, and BoringSSL and whether they are patched ✓ as of October 2019, currently being patched 🔧, or unpatched ✗. Exploiting the side channel can be easy ●, medium ①, or hard ○. The number of leaked bits (Nonce Leakage) indicates the complexity of a full key recovery.

Vulnerability		OpenSSL	LibreSSL	BoringSSL	Nonce Leakage	SC	Comments
Exp. / Scalar Mult.	Generate: (V1) Small k (top)	EC✗	EC✗	–	Topmost 0-limbs of k	●	Leaks in several subsequent steps
	(V2) k -padding resize	DSA✓EC✓	DSA✗EC✗	–	Topmost 0-bits of k	●	CVE-2018-0734 and CVE-2018-0735
	(V3) consttime-swap	DSA✓EC✓	DSA✗EC✗	–	same as (V2)	①	Already known
	(V4) Downgrade	DSA✓	–	–	same as (V2) + [24]	●	Introduced while fixing (V2)
	(V5) k -padding (top)	DSA✗EC✗	DSA✗EC✗	–	same as (V2)	○	Leaks in BN_add and BN_is_bit_set. SGX attack shown in Appendix B.
	(V6) Buffer conversion	EC✓	–	–	Topmost 0-bytes of k	○	
	(V7) Point addition	EC🔧	–	EC✓	All 0-windows of k	○	
Invert	(V8) Euclid BN_div	DSA✓	DSA✗	–	Topmost bit of k	●	Leaks via resize, similar to (V2)
	(V9) Euclid negation	DSA✓	DSA✗	–	Topmost 0-bit of $kinv$	●	Leaks via conditional negation
Multiply: (V10) Small k^{-1} (top)		–	EC✓	–	Topmost 0-limbs of $kinv$	①	

Weiser et al., “Big Numbers - Big Troubles: Systematically Analyzing Nonce Leakage in (EC)DSA Implementations”, USENIX 2020,
<https://www.usenix.org/system/files/sec20-weiser.pdf>

Breaking ECDSA with Partially Known Nonces

- You will implement a synthetic version of these attacks in the lab to recover signing keys.
- The remainder of these lectures will describe how you will do this, in a sequence of steps:
 - Reducing to Closest Vector Problem (CVP) in a lattice.
 - Reducing CVP to Shortest Vector Problem (SVP).
 - Solving SVP using the LLL algorithm (as a black box).
- The original ideas go back to: Howgrave-Graham and Smart, Lattice Attacks on Digital Signature Schemes. *Des. Codes Cryptography*, 23:283–290, 2001.

Lattices and Lattice Reduction

Lattices


Definition (full-rank lattice)

Let $\{\underline{b}_1, \underline{b}_2, \dots, \underline{b}_n\}$ be n linearly independent vectors in \mathbb{R}^n .

Then the lattice generated by $\{\underline{b}_1, \underline{b}_2, \dots, \underline{b}_n\}$ is the set:

$$L := \{ \sum_{i=1}^n l_i \underline{b}_i : l_i \in \mathbb{Z} \}$$

of integer linear combinations of the \underline{b}_i .

- If we allow real-valued linear combinations, we just get $L = \mathbb{R}^n$.
- Example: $n=2$; $\underline{b}_1 = (1,0)$; $\underline{b}_2 = (0,1)$, then $L = \mathbb{Z} \times \mathbb{Z}$.
- Example: $n=2$; $\underline{b}_1 = (12.3, \pi)$; $\underline{b}_2 = (-5, e)$, then $L = ??$.
- In general, the sums, differences etc, of lattice vectors are also lattice vectors. 
- So a lattice is a discrete subgroup of \mathbb{R}^n .

Lattice Bases

Definition (basis matrix of a full-rank lattice):

A basis matrix B of a lattice $L \subset \mathbb{R}^n$ is an $n \times n$ matrix formed by taking the rows to be basis vectors \underline{b}_i . Then

$$L = \{ \underline{x}B : \underline{x} \in \mathbb{Z}^n \}$$

- In cryptanalysis applications, our matrix entries will usually be integers.
- Two different, full-rank matrices B, B' can generate the same lattice.

Fact:

Two different matrices B, B' generate the same lattice L if and only if $B' = UB$ where U is an $n \times n$ matrix with integer entries and determinant ± 1 .

Definition (determinant of a full-rank lattice):

The determinant of a full-rank lattice L is the absolute value of the determinant of any basis matrix B for the lattice.

Lattice Successive Minima

Definition (norm of a vector):

For $\underline{v} \in \mathbb{R}^n$, $\|\underline{v}\|$ denotes the Euclidean norm of \underline{v} , i.e.

$$\|\underline{v}\| = (\sum_{i=1}^n v_i^2)^{1/2}.$$

Definition (successive minima of a lattice):

Let $L \subset \mathbb{R}^n$ be a full rank lattice. The successive minima of L are the values $\lambda_1, \dots, \lambda_n \in \mathbb{R}$ such that:

For $1 \leq i \leq n$, λ_i is the smallest real value such that there exist i linearly independent vectors $\underline{v}_1, \dots, \underline{v}_i$ with $\|\underline{v}_j\| \leq \lambda_i$ for $1 \leq j \leq i$.

Special case:

λ_1 is the length of a **shortest** (in terms of Euclidean norm) non-zero vector in L .

The Gaussian Heuristic

There are many bounds on the size of the successive minima, particularly for λ_1 .

The Gaussian Heuristic

Let $L \subset \mathbb{R}^n$ be a “random” full rank lattice. Then

$$\lambda_1 \approx (n/2\pi e)^{1/2} \cdot \det(L)^{1/n}.$$

NB: “*random lattice*” is not formally defined, but the ones we use in cryptanalysis can often be assumed to behave in this way.

SVP and CVP

Let $L \subset \mathbb{R}^n$ be a “random” full rank lattice, and let $\underline{w} \in \mathbb{R}^n$.

Then the **Shortest Vector Problem (SVP)** is to find $\underline{v} \in L$ such that $\|\underline{v}\| = \lambda_1$.

The **Closest Vector Problem (CVP)** for \underline{w} is to find $\underline{v} \in L$ such that $\|\underline{v} - \underline{w}\|$ is as small as possible.

- SVP and CVP are known to be hard problems in general.
- For example, CVP can be shown to be NP-hard by relating it to subset-sum.
- SVP and CVP problems in high-dimensional lattices can be used to **construct** cryptographic schemes, e.g. public key encryption, signatures.
- However, SVP and CVP may be easy if the dimension is small and/or the lattice is presented in a “nice” way.
- Then lattices can be used as a tool for **cryptanalysis**.

Lattice Reduction

- Let $L \subset \mathbb{R}^n$ be a full rank lattice, represented by some basis matrix B .
- **Lattice reduction** refers to the process of producing a new basis matrix B' for L satisfying certain special properties.
- In particular, the rows of B' (whose linear combinations define L) are “somewhat orthogonal” and the norms of the first rows of B' are relatively short.
- The **Lenstra, Lenstra, Lovasz (LLL) algorithm** is a deterministic algorithm which performs lattice reduction.
- LLL essentially performs “iterative rounded Gram-Schmidt orthogonalization”.
- Algorithmic details can be found in Chapter 17 of Galbraith’s book “Mathematics of Public Key Cryptography” available for free at:
<https://www.math.auckland.ac.nz/~sgalo18/crypto-book/crypto-book.html>

Lattice Reduction – LLL and BKZ


- LLL runs in time (and space) polynomial in n and $\max_i \|\underline{b}_i\|^2$ and produces a basis matrix B' in which the first row \underline{b}'_1 satisfies:

$$\|\underline{b}'_1\| \leq 2^{(n-1)/2} \lambda_1.$$

- In other words, LLL *approximately* solves the SVP for lattice L (and more).
- In practice, LLL often *exactly* solves SVP.
- For large dimensions n , LLL is superseded by the BKZ algorithm (which uses LLL as a subroutine to solve SVP on sub-lattices).
- Due to its importance in cryptography, lattice reduction is a major area of on-going research, with many improvements to BKZ in recent years.
- In the lab, we will be using **fpyl**, a state-of-the-art lattice reduction package for Python.
- Nice intro at: <https://martinralbrecht.wordpress.com/2016/04/03/fpylll/>

From ECDSA to CVP

From ECDSA to CVP

- Suppose we are given a signature (r, s) on message m .
- Suppose also that the L MSBs of the corresponding nonce k are known.
 - Let N be the bit-length of q . 
 - Then k lies in the interval $[a2^{N-L}, (a+1)2^{N-L}-1]$ for some known a (coming from MSBs of k).
 - The mid-point of this interval is $a2^{N-L} + 2^{N-L-1}$.
 - So let's write $k = a2^{N-L} + 2^{N-L-1} + e$ where $0 \leq |e| \leq 2^{N-L-1}$.
- Rearranging the signing equation $s = k^{-1}(h + xr) \bmod q$ yields:
$$(rs^{-1})x = k - hs^{-1} \bmod q$$
- Set $t = rs^{-1} \bmod q$ and $z = hs^{-1} \bmod q$ (z is an integer between 0 and $q-1$).
- Then we have:

$$tx = k - z \bmod q.$$

From ECDSA to CVP

- We have: $tx = k - z \bmod q.$

- Now:

$$k = a2^{N-L} + 2^{N-L-1} + e,$$

so

$$k - z = u + e$$

where

$$u = a2^{N-L} + 2^{N-L-1} - z$$

is a known integer that can be computed from the L MSBs of k and $z = hs^{-1} \bmod q$.

- Moreover, e is bounded by: $0 \leq |e| \leq 2^{N-L-1}.$
- So we finally arrive at: $tx = u + e \bmod q$, and hence:

$$tx = u + e + l \cdot q \text{ for some } l$$

- Since $u = a2^{N-L} + 2^{N-L-1} - z$, we see that u lies between $2^{N-L-1} - q$ and $a2^{N-L} + 2^{N-L-1}.$

- We really only care about values mod q , so we can assume (by adding multiples of q as needed) that u is centred, i.e. $-q/2 < u < q/2.$

From ECDSA to CVP

- So far: from signature (r, s) on message m and L MSBs of k we get:

$$tx = u + e + l \cdot q \quad \text{for some } l.$$

- Here, x is our target, t is known, u is known, e is small but otherwise unknown and l is unknown.
- We get one such equation for each of n signatures (r_i, s_i) on messages m_i :

$$t_i x = u_i + e_i + l_i \cdot q$$

- Alternatively, we can write:

$$t_i x = u_i + e_i \pmod{q} \quad (\text{where } e_i \text{ is small})$$

meaning that u_i is a good approximation to $t_i x \pmod{q}$.

- In this second form, the problem of recovering x from n distinct equations with uniformly random t_i is called the **Hidden Number Problem (HNP)**.
- A similar translation can be done when the **LSBs of the k_i are known, or in fact any set of contiguous bits.**

From ECDSA to CVP

A Formal Result:

Let t_1, \dots, t_n be uniformly random values in F_q , let x be non-zero in F_q . Suppose we are given n samples of the form (t_i, u_i) where u_i is known to be a good approximation to $t_i x \bmod q$ (i.e. $t_i x = u_i + e_i \bmod q$ with $0 \leq |e_i| \leq q/2^{L+1}$).

Suppose $n = 2\log_2(q)^{1/2}$ and $L = \log_2(q)^{1/2} + \log_2 \log_2(q)$.

Then one can recover x in polynomial time.

- For a proof, see Theorem 21.7.9 and Corollary 21.7.10 of Galbraith's book.
- The proof is based on properties of LLL and the Babai nearest plane algorithm for solving CVP.
- Guarantee required of $|e_i|$ is slightly stronger here than in our formulation.
- In practice, we can get away with much smaller n and L and still get an attack that works.

From ECDSA to CVP

Consider the lattice $L \subset \mathbb{R}^{n+1}$ with basis matrix B given by:

$$\begin{vmatrix} q & o & o & \dots & o & o \\ o & q & o & \dots & o & o \\ . & & & & . & . \\ . & . & . & \dots & . & . \\ . & & & & & \\ o & o & o & \dots & q & o \\ t_1 & t_2 & t_3 & & t_n & 1/2^{L+1} \end{vmatrix}$$

Define $\underline{u} = (u_1, u_2, \dots, u_n, o) \in \mathbb{R}^{n+1}$ where, recall, $t_i x = u_i + e_i + l_i \cdot q$ with e_i small.

Claim: There exists a vector $\underline{v} \in L$ such that

$$\|\underline{u} - \underline{v}\| < (n+1)^{1/2} \cdot 2^{N-L-1}.$$

From ECDSA to CVP

Proof of claim:

We can write

$$t_i x = u_i + e_i + l_i \cdot q$$

where $|e_i| \leq 2^{N-L-1}$, and l_i is some (unknown) integer.

Now define $\underline{v} \in L$ via:

$$\underline{v} = (-l_1, -l_2, \dots, -l_n, x) \cdot B.$$

(Recall that x is unknown, so \underline{v} is also unknown at this point.)

$$\begin{aligned} \underline{v} &= (-l_1, -l_2, \dots, -l_n, x) \cdot B = (-l_1 q + t_1 x, -l_2 q + t_2 x, \dots, -l_n q + t_n x, x/2^{L+1}) \\ &= (u_1 + e_1, u_2 + e_2, \dots, u_n + e_n, x/2^{L+1}) \end{aligned}$$

Hence $\underline{v} - \underline{u} = (e_1, e_2, \dots, e_n, x/2^{L+1})$ and the result follows on noting that each entry in $\underline{v} - \underline{u}$ is bounded in absolute value by 2^{N-L-1} (for the last entry, note that $x < q < 2^N$).

From ECDSA to CVP

Implications:

- We have constructed a lattice $L \subset \mathbb{R}^{n+1}$ and vector \underline{u} from public information and shown that \underline{u} is “somewhat close” to a point \underline{v} in L .
- Moreover, finding \underline{v} allows us to find the ECDSA private key x (just by inspecting the final coordinate of \underline{v}).
- We can hope that \underline{v} is actually **the** solution to the CVP for \underline{u} .
- For some parameter ranges, one can show that this is indeed the case.
- So if we have a CVP solver, we can apply it here and hope to extract the private key x .
- So how do we solve CVP?

From CVP to SVP: Kannan's Embedding Technique

From CVP to SVP: Kannan's Embedding Technique

- There are multiple ways to solve CVP using an SVP solver: the Babai nearest plane algorithm, Babai rounding, Kannan's embedding technique, enumeration approaches.
- We will describe only Kannan embedding here, as it is nice and simple.
- Babai nearest plane and Babai rounding are also simple, and have provable guarantees.
- Your generic LLL library may allow you to solve CVP directly, but it's good to have a sense of what could be happening underneath!

From CVP to SVP: Kannan's Embedding Technique

Let B be a basis matrix for a lattice $L \subset \mathbb{R}^n$ with rows \underline{b}_i .



Let $\underline{w} \in \mathbb{R}^n$ be a vector for which we wish to solve CVP.

A solution to the CVP corresponds to integers l_1, \dots, l_n such that:

$$\underline{w} \approx l_1 \underline{b}_1 + \dots + l_n \underline{b}_n.$$

Define $\underline{f} = \underline{w} - (l_1 \underline{b}_1 + \dots + l_n \underline{b}_n)$.



Key observation: $\|\underline{f}\|$ is small.

So we try to define a new lattice L' which contains \underline{f} – hopefully then \underline{f} will be output as a result of running an SVP solver on L' .

From CVP to SVP: Kannan's Embedding Technique

- Consider the lattice $L' \subset \mathbb{R}^{n+1}$ with basis matrix B' whose rows are:

$$(\underline{b}_1, 0), (\underline{b}_2, 0), \dots, (\underline{b}_n, 0), (\underline{w}, M)$$

where, recall, \underline{w} is the input to the CVP.

- Here M is a constant, to be determined.
- Now consider the linear combination of rows with coefficients:

$$(-l_1, \dots, -l_n, 1)$$

- It is easy to check that this yields the vector (\underline{f}, M) , which should be short.
- So we might be able to solve CVP on input \underline{w} for lattice L by solving SVP on lattice L' to find (\underline{f}, M) and then setting $\underline{v} = \underline{w} - \underline{f}$.

From CVP to SVP: Kannan's Embedding Technique

Lemma:

Let $L \subset \mathbb{R}^n$ be a full rank lattice with shortest non-zero vector of length λ_1 . Let $\underline{w} \in \mathbb{R}^n$ and let \underline{v} be a closest lattice vector to \underline{w} . Define $\underline{f} = \underline{w} - \underline{v}$. Suppose that $\|\underline{f}\| \leq \lambda_1/2$ and let $M = \|\underline{f}\|$. Then (\underline{f}, M) is a shortest vector in the lattice $L' \subset \mathbb{R}^{n+1}$ in Kannan's embedding technique.

Proof: see Galbraith's book, Lemma 18.3.2.

Interpretation: if the target vector \underline{w} is very close to the lattice and we have a good guess M for the distance, then Kannan's embedding technique does reduce the problem of solving CVP to that of solving SVP.

Problems: maybe \underline{w} is not close to the lattice; LLL and related algorithms only approximately solve SVP; maybe target (\underline{f}, M) is short but **not** a shortest vector in L' .

Solution: in practice, this approach works quite well, but we may need to examine several vectors in the reduced basis to find target (\underline{f}, M) or perform *enumeration*.

Putting It All Together

Putting It All Together

- We have seen how to translate the problem of recovering x from partial information about the nonces in the ECDSA scheme into a CVP problem and thence to an SVP problem.
- The lattice dimension we use is $n+2$ where n is the number of signatures.
- Whether n signatures with L bits of leakage per signature is enough to recover k depends on several factors.
- Clearly there is an information theoretic minimum: we need

$$n \cdot L > \log_2(q) = N.$$

- Actually, knowing the public key $[x]P$ enables attack to go beyond this minimum.

Putting It All Together

- We can use the Gaussian heuristic to see if Kannan's embedding technique or some other CVP solver is likely to produce a solution – see the exercises for an important wrinkle in this approach.
- You will implement all this in the lab and use fpylll as a tool to solve the lattice instances that arise.
- Through this programming exercise, you'll explore the performance of this kind of attack and deal with some of the subtleties that arise.
- Have fun!